

# Host Compiled Simulation for Timing and Power Estimation

Gaurav Kukreja  
Bo Wang

Intel Mobile Communications

October 29, 2014

# Overview

Introduction

Simple Example

Timing Estimation

Power Estimation

Conclusion

# Simulation

Simulation is the technique to imitate the behaviour of a system.

- ▶ Widely used in Hardware Software Co-development.
- ▶ Use cases are performance analysis, functional verification etc.

# Simulation: Popular Techniques

## Cycle Accurate Simulation

- ▶ Detailed simulation of processor micro-architecture.
- ▶ Cycle Accurate estimation of performance.
- ▶ Difficult to develop, Very Slow execution.

## Functional Simulation

- ▶ High Level of Abstraction.
- ▶ Simple to develop, and fast execution.
- ▶ Focus is Functional Verification. Cannot be used for performance estimation.

# Our Focus

A technique for fast simulation of embedded processors that is,

- ▶ Easy to Develop.
- ▶ Fast to Execute.
- ▶ Highly Accurate in Performance Estimation.

## Related Work

### Sampling Based Approach

- ▶ Small Samples Executed using CAS.
- ▶ Rest execution fast-forwarded using Functional Simulation.
- ▶ Results Interpolated.
- ▶ Inaccurate, Difficult to develop (CAS)

# Host Compiled Simulation

## Host Compiled Simulation

- ▶ Based on technique of Source Code Instrumentation.
- ▶ Instrumented code compiled and run on Host Machine, hence the name.
- ▶ Easy to understand, develop and maintain.
- ▶ Fast Execution, and accurate results.

## Simple Example

```

1  int sum(int array[20])
2  {
3      int i;
4      int sum = 0;
5
6      for (i=0; i<20; i++)
7          sum += array[i];
8
9      return sum;
10 }
```

Simple C Code

```

1  00008068 <sum>:
2  8068:  mov     r3, #0
3  806c:  mov     r2, r3
4  8070:  ldr     r1, [r0, r3]
5  8074:  add     r2, r2, r1
6  8078:  add     r3, r3, #4
7  807c:  cmp     r3, #80 ; 0x50
8  8080:  bne     8070 <sum+0x8>
9  8084:  mov     r0, r2
10 8088:  bx      lr
```

Objdump Code



## Simple Example

```

1  int sum(int array[20])
2  {
3      int i;
4      int sum = 0;
5
6      for (i=0; i<20; i++)
7          sum += array[i];
8
9      return sum;
10 }
```

Simple C Code

```

1  00008068 <sum>:
2  8068:    mov     r3, #0
3  806c:    mov     r2, r3
4  8070:    ldr     r1, [r0, r3]
5  8074:    add     r2, r2, r1
6  8078:    add     r3, r3, #4
7  807c:    cmp     r3, #80 ; 0x50
8  8080:    bne     8070 <sum+0x8>
9  8084:    mov     r0, r2
10 8088:    bx      lr
```

Objdump Code

Basic Block in Binary		Matching block in Source	
BlockID	Lines	BlockID	Lines
1	2-3	1	3-4
2	4-8	2	7
3	9-10	3	9

# Instrumented Code

```

1  unsigned int execCycles;
2  unsigned int memAccessCycles;
3
4  int sum(int array[20])
5  {
6      int i;
7      int sum = 0;
8      execCycles += 2;
9      memAccessCycles += simICache(0x8068, 8);
10
11     for (i=0; i<20; i++)
12     {
13         sum += array[i];
14         memAccessCycles += simDCache(array + i, READ);
15         execCycles += 5;
16         memAccessCycles += simICache(0x8070, 40);
17     }
18
19     execCycles += 2;
20     memAccessCycles += simICache(0x8084, 8);
21     return sum;
22 }

```

```

1  int sum(int array[20])
2  {
3      int i;
4      int sum = 0;
5
6      for (i=0; i<20; i++)
7          sum += array[i];
8
9      return sum;
10 }

```

```

1  00008068 <sum>:
2  8068:  mov     r3, #0
3  806c:  mov     r2, r3
4  8070:  ldr     r1, [r0, r3]
5  8074:  add     r2, r2, r1
6  8078:  add     r3, r3, #4
7  807c:  cmp     r3, #80 ; 0x50
8  8080:  bne     8070 <sum+0x8>
9  8084:  mov     r0, r2
10 8088:  bx      lr

```

# Objective

- ▶ Develop a tool for Automatic Instrumentation.
- ▶ ARM Cortex A5 based processor as reference target device.
- ▶ Bare-Metal Applications.
- ▶ Generate Time and Power Consumption Estimates.

## Mapping between Source and Binary

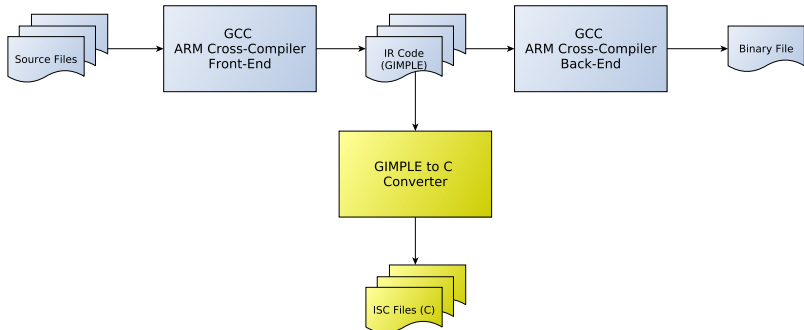
- ▶ Very important for accurate instrumentation.
- ▶ Compiler destroys mapping during optimization phases.
- ▶ GDB provides mapping, but highly inaccurate.
- ▶ Control Flow Analysis to generate mapping between Basic Blocks.

## Mapping between Source and Binary

In this project, mapping is generated using following steps.

- ▶ Cross-Compile Source Code.
- ▶ Convert IR Code to C Code (Intermediate Source Code).
- ▶ Extract CFG from ISC and Binary Code.
- ▶ Map CFGs using Mapping Algorithm

## Conversion of IR Code to C Code



\* Code for GIMPLE to C converter reused from RBA project [TODO]

# Mapping Algorithm

- ▶ Extract CFG from ISC and Binary Code.
- ▶ Graph Matching Algorithm using Depth First Traversal.
- ▶ Special handling for each optimization.
- ▶ GDB Debug information in Corner Cases.


# Handling of Compiler Optimization : Conditional Execution

```

1  ...
2  if (a > b)
3      max = a;
4  else
5      max = b;
6  ...

```

Simple C Code




```

1  806c:    ...
2  8070:    cmp     r1, r2
3  8074:    ble     8080
4  8078:    mov     r3, r1
5  807c:    b        8084
6  8080:    mov     r3, r2
7  8084:    ...

```

Unoptimized Assembly Code





# Handling of Compiler Optimization : Conditional Execution

```

1  ...
2  if (a > b)
3      max = a;
4  else
5      max = b;
6  ...

```

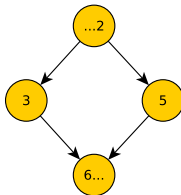
Simple C Code

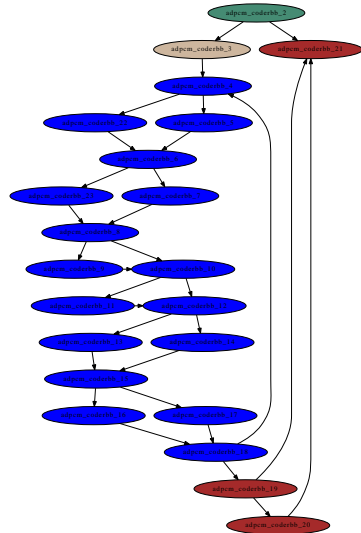

```

1  806c:  ...
2  8070:  cmp    r1, r2
3  8074:  movgt  r3, r1
4  8078:  movle  r3, r2
5  807c:  ...

```

Optimized Assembly Code





## Annotation for Cycles spent in Pipeline

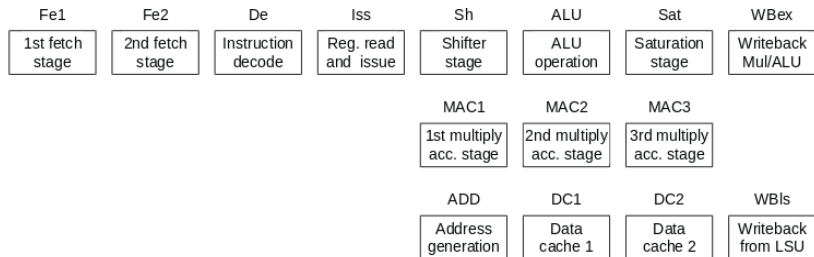
Performed at Basic Block Granularity.

Important to consider

- ▶ Pipeline Architecture of the target processor.
- ▶ Data and Control Hazards, that leads to pipeline stalls.
- ▶ Branch Prediction, that prevents pipeline flushes.

## Pipeline architecture of ARM Cortex A5

The ARM Cortex A5 has an 8-stage pipeline.



Lets assume,

- ▶ Pipeline is empty at start of Basic Block.
- ▶ All data is available without latency.

## Effects due to Data and Control Hazards

For each Basic Block from Binary Code,

- ▶ Parse instructions sequentially.
- ▶ Simulate progress in pipeline.
- ▶ Identify interlocking (hazards) and add penalty.
- ▶ Annotate to mapped block in Source Code.

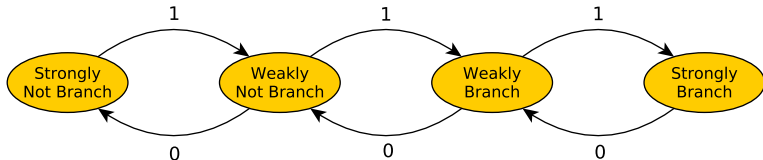
```
...  
for(i<0; i<10; i++) {}  
    execCycles += 23;  
    sum += array[i];  
    ...  
}  
...
```

## Branch Prediction

- ▶ Reduces number of pipeline flushes.
- ▶ Major impact on performance.
- ▶ Branch Prediction Unit is simulated to account for this.

## Branch Prediction Algorithm

- ▶ Simulates the BPU on ARM Cortex A5
- ▶ 125 entry Branch History Table.
- ▶ State Machine for each Entry. 2 bit state information.



# Branch Prediction API

Branch Prediction Simulator offers following API.

```
/**  
 * @brief Function called at the beginning of a Basic Block  
 *        to simulate Branch Prediction Unit.  
 *  
 * @param Start Address of Basic Block being entered.  
 * @param End Address of Basic Block being entered.  
 *  
 * @return True, if branch was correctly predicted.  
 *         False, if branch was not correctly predicted.  
 */  
unsigned int branchPred_enter(unsigned long startAdd,  
                               unsigned long endAdd);
```



## Annotation for Branch Prediction

- ▶ Start and End address of Basic Block extracted from Binary.
- ▶ Annotation done at beginning of Basic Block.
- ▶ Depending on outcome, penalty is subtracted from `execCycles`.

```
...  
for(i=0; i<10; i++) {  
    execCycles += 23;  
    execCycles -= (branchPred_enter(0x348, 0x380) ? 7 : 0);  
    sum += array[i];  
    ...  
}  
...
```

## Cache Simulator

Cache Hierarchy on target device.

	Size	N-way	Cache Line Size
L1 D Cache	32 KB	4	32 B
L1 I Cache	32 KB	2	32 B
L2 Cache	256 KB	16	32 B

- ▶ Pseudo Random Replacement Policy
- ▶ Data Prefetching

# Instruction Cache Simulation

```
/**  
 * @brief Function to simulate Instruction Cache Access.  
 *  
 * @param Start Address of the basic block.  
 * @param Size of the basic block in Bytes.  
 *  
 * @return Number of cycles spent in performing access.  
 */  
unsigned long long simICache(unsigned long address,  
                             unsigned long size);
```

## Instruction Cache Simulation

- ▶ Start address and size of Basic Blocks extracted from Binary.
- ▶ Annotation at beginning of mapped Basic Block in Source Code.
- ▶ Return value added to `memAccessCycles`.

```

1  ...
2  for(i=0; i<10; i++) {
3      execCycles += 23;
4      execCycles -= (branchPred_enter(0x348, 0x380) ? 7 : 0);
5      memAccessCycles += simICache(0x348, 56);
6      sum += array[i];
7      ...
8  }
9  ...

```

## Data Cache Simulation

```
/**  
 * @brief Function to simulate Data Cache Access.  
 *  
 * @param Address of the memory being accessed.  
 * @param True, if access is Read Access.  
 *        False, if write access.  
 *  
 * @return Number of cycles spent in performing access.  
 */  
unsigned long long simDCache(unsigned long address,  
                             unsigned int isReadAccess);
```

## Data Cache Simulation

Address from host cannot be used for simulation.

Will lead to inaccuracies, because

- ▶ Different sizes of Basic Data Types.
- ▶ Different Memory Alignment on Host and Target.

Memory Access Reconstruction

- ▶ Inspired by [TODO]
- ▶ Resolve address of each load/store instruction on target device.

# Memory Access Reconstruction

- ▶ Resolve Physical Address of each variable in target memory.
- ▶ Identify variable being accessed by each load/store.
- ▶ Identify index information from source code.
- ▶ Simulate Data Cache.

# Resolve Physical Address of Variables

## Global Variables

- ▶ Bare-Metal Execution. Address fixed at compile time.
- ▶ Extracted by Static Analysis of binary.

Annotate address in code as follows,

```
1 unsigned int globVar[100];  
2 unsigned long globVar_addr = 0x7c8;
```



## Resolve Physical Address of Variables

### Local Variable

- ▶ Stored in Stack, address only known at run-time.
- ▶ Address relative to stack, can be known statically.
- ▶ Simulate growth of stack at run-time to resolve physical address.

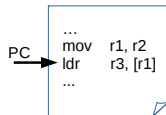
```

1  unsigned long CSIM_SP = 0x1d4c8;           // Initial Value of SP.
2
3  int foo()
4  {
5      CSIM_SP += 88;                         // Size of Stack Frame for foo.
6
7      int localVar;
8      unsigned long localVar_addr = 0x8;     // Address relative to SP.
9      ...
10 }
```

## Identify Load/Store Operations in Binary Code

To identify variable accessed by a load/store instruction,

- ▶ Binary code is functionally simulated.
- ▶ Register state is maintained, and updated as per instruction.
- ▶ Branch instructions are ignored.
- ▶ Address for each load/store is extracted.
- ▶ Variable being accessed is known.



Assembly Code

r0	r1	r2	...	r15
0	7c8	7c8	...	346

Register State

## Annotation of Memory Access

To find the line in source code, which causes the memory access operation,

- ▶ Each line is parsed using a custom C Parser.
- ▶ Index information is extracted.
- ▶ Annotation as follows.

```

1  ...
2  for (i=0; i < 100; i++) {
3      sum += globalVar[i];
4      memAccessCycles += simDCache ( globalVar_addr + i * 4, True );
5  }
6  localVar = sum / 100;
7  memAccessCycles += simDCache ( SP + localVar_addr, False );
8  ...

```


## Timing Estimation

- ▶ Annotated Source Code is compiled and run on Host Machine.
- ▶ Time spent in Active state and Idle State of CPU is reported.
- ▶ Additionally, trace information from Cache is generated.

For estimating power, the power state model approach has been used.

- ▶ Power consumed by each component in active and idle state is known.
- ▶ Trace information for each basic block, shows components that were active and the duration.
- ▶ Total power being consumed at a time, can be estimated.

## Components for Power Estimation



Target System is divided into components as follows

- ▶ Core and L1 Caches
- ▶ L2 Cache
- ▶ External Memory along with controller and NOC.

## Test Setup


For testing, ARM Cortex A5 core was used. Results from ARM Performance Measurement Unit were compared.

- ▶ Number of Cache Misses.
- ▶ Total Cycles.

Accurate Prediction of Cache Misses, will verify the correctness of the instrumentation.

# ADPCM Benchmark

	HCS	Actual	Accuracy
Total Cache Miss	91452	91604	99.99%
Total Cycles	46110394	45594862	99.98%





# Sieve Benchmark

	HCS	Actual	Accuracy
Total Cache Miss	933109	934250	99.88%
Total Cycles	99068687	100403731	98.67%

