

Host Compiled Simulation for Timing and Power Estimation

Gaurav Kukreja
Bo Wang

Intel Mobile Communications

October 29, 2014

Overview

Introduction

Simple Example

Timing Estimation

Power Estimation

Conclusion

Simulation: Popular Techniques

Cycle Accurate Simulation

- ▶ Detailed simulation of processor micro-architecture.
- ▶ Cycle Accurate estimation of performance.
- ▶ Difficult to develop, Very Slow execution.

Functional Simulation

- ▶ High Level of Abstraction.
- ▶ Simple to develop, and fast execution.
- ▶ Focus is Functional Verification. Cannot be used for performance estimation.

Our Focus

A technique that is,

- ▶ Easy to Develop.
- ▶ Fast to Execute.
- ▶ Accurate.

Host Compiled Simulation

Host Compiled Simulation

- ▶ Based on Source Code Instrumentation.
- ▶ Instrumented code compiled and run on Host Machine, hence the name.
- ▶ Easy to understand, develop and maintain.
- ▶ Fast Execution, and accurate results.

Simple Example

```
1  int sum(int array[20])
2  {
3      int i;
4      int sum = 0;
5
6      for (i=0; i<20; i++)
7          sum += array[i];
8
9      return sum;
10 }
```

Simple C Code

```
1  00008068 <sum>:
2  8068:    mov     r3, #0
3  806c:    mov     r2, r3
4  8070:    ldr     r1, [r0, r3]
5  8074:    add     r2, r2, r1
6  8078:    add     r3, r3, #4
7  807c:    cmp     r3, #80 ; 0x50
8  8080:    bne     8070 <sum+0x8>
9  8084:    mov     r0, r2
10 8088:    bx      lr
```

Objdump Code

Simple Example: Mapping Basic Blocks

```

1  int sum(int array[20])
2  {
3      int i;
4      int sum = 0;
5
6      for (i=0; i<20; i++)
7          sum += array[i];
8
9      return sum;
10 }
```

Simple C Code

```

1  00008068 <sum>:
2  8068:    mov     r3, #0
3  806c:    mov     r2, r3
4  8070:    ldr     r1, [r0, r3]
5  8074:    add     r2, r2, r1
6  8078:    add     r3, r3, #4
7  807c:    cmp     r3, #80 ; 0x50
8  8080:    bne     8070 <sum+0x8>
9  8084:    mov     r0, r2
10 8088:    bx      lr
```

Objdump Code

Basic Block in Binary		Matching block in Source	
BlockID	Lines	BlockID	Lines
1	2-3	1	3-4
2	4-8	2	6-7
3	9-10	3	9

Instrumented Code

```

1  unsigned int execCycles;
2  unsigned int memAccessCycles;
3
4  int sum(int array[20])
5  {
6      int i;
7      int sum = 0;
8      execCycles += 2;
9      memAccessCycles += simICache(0x8068, 8);
10
11     for (i=0; i<20; i++)
12     {
13         sum += array[i];
14         memAccessCycles += simDCache(array + i, READ);
15         execCycles += 5;
16         memAccessCycles += simICache(0x8070, 40);
17     }
18
19     execCycles += 2;
20     memAccessCycles += simICache(0x8084, 8);
21     return sum;
22 }

```

```

1  int sum(int array[20])
2  {
3      int i;
4      int sum = 0;
5
6      for (i=0; i<20; i++)
7          sum += array[i];
8
9      return sum;
10 }

```

```

1  00008068 <sum>:
2  8068:  mov     r3, #0
3  806c:  mov     r2, r3
4  8070:  ldr     r1, [r0, r3]
5  8074:  add     r2, r2, r1
6  8078:  add     r3, r3, #4
7  807c:  cmp     r3, #80 ; 0x50
8  8080:  bne     8070 <sum+0x8>
9  8084:  mov     r0, r2
10 8088:  bx      lr

```


Objective: Develop a tool for Automatic Instrumentation

Steps:

- ▶ Extract Mapping.
- ▶ Instrumentation for cycles spent in Execution.
- ▶ Instrumentation for cycles spent in Memory Access.

Output:

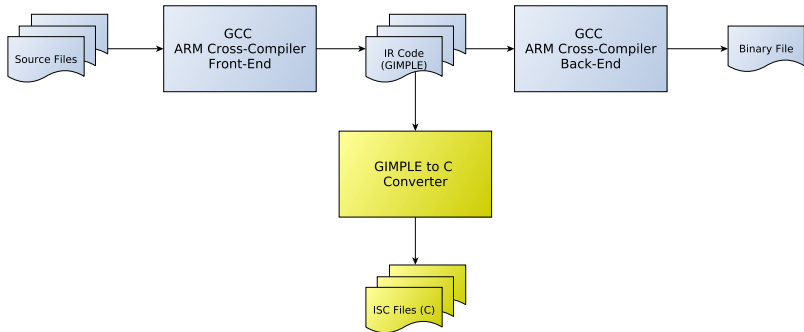
- ▶ Instrumented Code.

Extract Mapping between Source Code and Binary Code

Mapping between Source and Binary

- ▶ Very important for accurate instrumentation.
- ▶ Compiler destroys mapping during optimization phases.
- ▶ GDB provides mapping, but ambiguous.
- ▶ Control Flow Analysis to generate mapping between Basic Blocks.

Conversion of IR Code to C Code



* Code for GIMPLE to C converter reused from RBA project [1]

Mapping Algorithm

- ▶ Extract CFG from ISC and Binary Code.
- ▶ Graph Matching Algorithm using Depth First Traversal.
- ▶ Special handling for each optimization.
- ▶ GDB debug information in corner cases.


Handling of Compiler Optimization : Conditional Execution

```

1  ...
2  if (a > b)
3      max = a;
4  else
5      max = b;
6  ...

```

Simple C Code




```

1  806c:  ...
2  8070:  cmp     r1, r2
3  8074:  ble     8080
4  8078:  mov     r3, r1
5  807c:  b        8084
6  8080:  mov     r3, r2
7  8084:  ...

```

Unoptimized Assembly Code




Handling of Compiler Optimization : Conditional Execution



```
1  ...  
2  if (a > b)  
3      max = a;  
4  else  
5      max = b;  
6  ...
```



Simple C Code

```
1  806c:  ...  
2  8070:  cmp    r1, r2  
3  8074:  movgt  r3, r1  
4  8078:  movle  r3, r2  
5  807c:  ...
```

Optimized Assembly Code







Annotation for Cycles Spent in Execution

Annotation for Cycles spent in Pipeline

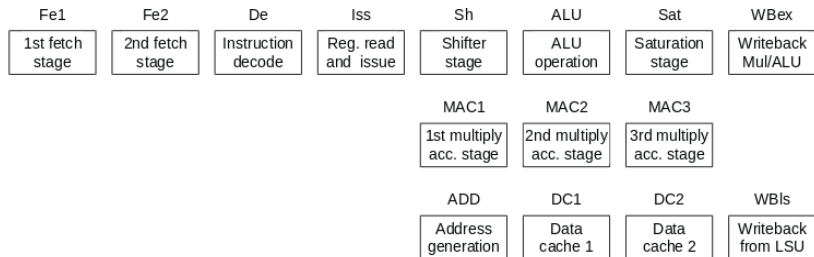
Performed at Basic Block Granularity.

Important to consider

- ▶ Pipeline Architecture of the target processor.
- ▶ Data and Control Hazards, that leads to pipeline stalls.
- ▶ Branch Prediction, that prevents pipeline flushes.

Pipeline architecture of ARM Cortex A5

The ARM Cortex A5 has an 8-stage pipeline.



Lets assume,

- ▶ Pipeline is empty at start of Basic Block.
- ▶ All data is available without latency.

Effects due to Data and Control Hazards

For each Basic Block from Binary Code,

- ▶ Parse instructions sequentially.
- ▶ Simulate progress in pipeline.
- ▶ Identify interlocking (hazards) and add penalty.
- ▶ Annotate to mapped block in Source Code.

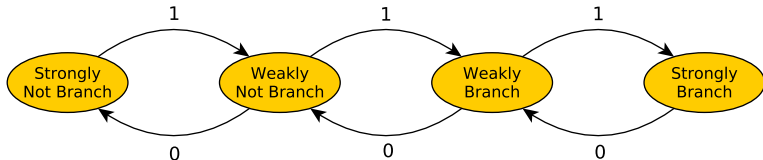
```
...  
for(i<0; i<10; i++) {}  
    execCycles += 23;  
    sum += array[i];  
    ...  
}  
...
```

Branch Prediction

- ▶ Reduces number of pipeline flushes.
- ▶ Major impact on performance.
- ▶ Branch Prediction Unit is simulated to account for this.

Branch Prediction Algorithm

- ▶ Simulates the BPU on ARM Cortex A5
- ▶ 125 entry Branch History Table.
- ▶ State Machine for each Entry. 2 bit state information.



Annotation for Branch Prediction

Branch Prediction Simulator offers following API.

```
unsigned int branchPred_enter(unsigned long startAdd,
                              unsigned long endAdd);
```

```
1  ...
2  for(i=0; i<10; i++) {
3      execCycles += 23;
4      execCycles -= (branchPred_enter(0x348, 0x380) ? 7 : 0);
5      sum += array[i];
6      ...
7  }
8  ...
```


Annotation for Cycles spent in Memory Access

Cache Simulator

Cache Hierarchy on target device.

	Size	N-way	Cache Line Size
L1 D Cache	32 KB	4	32 B
L1 I Cache	32 KB	2	32 B
L2 Cache	256 KB	16	32 B

- ▶ Pseudo Random Replacement Policy
- ▶ Data Prefetching

Instruction Cache Simulation

Cache Simulator offers API for Instruction Access.

```
unsigned long long simICache(unsigned long address,
                             unsigned long size);
```

```
1  ...
2  for(i=0; i<10; i++) {
3      execCycles += 23;
4      execCycles -= (branchPred_enter(0x348, 0x380) ? 7 : 0);
5      memAccessCycles += simICache(0x348, 56);
6      sum += array[i];
7      ...
8  }
9  ...
```

Data Cache Simulation

Address from host cannot be used for simulation.

Will lead to inaccuracies, because

- ▶ Different sizes of Basic Data Types.
- ▶ Different Memory Alignment on Host and Target.

Memory Access Reconstruction

- ▶ Inspired by [2]
- ▶ Resolve address of each load/store instruction on target device.

Memory Access Reconstruction

- ▶ Resolve Physical Address of each variable in target memory.
- ▶ Identify variable being accessed by each load/store.
- ▶ Identify index information from source code.
- ▶ Simulate Data Cache.

Resolve Physical Address of Variables

Global Variables

- ▶ Bare-Metal Execution. Address fixed at compile time.
- ▶ Extracted by Static Analysis of binary.

Annotate address in code as follows,

```
1 unsigned int globVar[100];  
2 unsigned long globVar_addr = 0x7c8;
```

Resolve Physical Address of Variables

Local Variable

- ▶ Stored in Stack, address only known at run-time.
- ▶ Address relative to stack, can be known statically.
- ▶ Simulate growth of stack at run-time to resolve physical address.


```

1  unsigned long CSIM_SP = 0x1d4c8;           // Initial Value of SP.
2
3  int foo()
4  {
5      CSIM_SP += 88;                         // Size of Stack Frame for foo.
6
7      int localVar;
8      unsigned long localVar_addr = 0x8;     // Address relative to SP.
9      ...
10 }
```

Identify Load/Store Operations in Binary Code

To identify variable accessed by a load/store instruction,

- ▶ Binary code is functionally simulated.
- ▶ Register state is maintained, and updated as per instruction.
- ▶ Branch instructions are ignored.
- ▶ Address for each load/store is extracted.
- ▶ Variable being accessed is known.



Assembly Code

r0	r1	r2	...	r15
0	7c8	7c8	...	346

Register State

Annotation of Memory Access

To find the line in source code, which causes the memory access operation,

- ▶ Each line is parsed using a custom C Parser.
- ▶ Index information is extracted.
- ▶ Annotation as follows.

```

1  ...
2  for (i=0; i < 100; i++) {
3      sum += globalVar[i];
4      memAccessCycles += simDCache ( globalVar_addr + i * 4, True );
5  }
6  localVar = sum / 100;
7  memAccessCycles += simDCache ( SP + localVar_addr, False );
8  ...


```

Power Estimation

Power State Model

- ▶ Power consumed by each component in active and idle state is known.
- ▶ Trace information for each basic block, shows components that were active and the duration.
- ▶ Total power being consumed at a time, can be estimated.

Components for Power Estimation



Target System is divided into components as follows

- ▶ Core and L1 Caches
- ▶ L2 Cache
- ▶ External Memory along with controller and NOC.

Limitaitons

- ▶ Instrumentation may not succeed in corner cases. Appropriate hints provided in source code as comments to manually perform instrumentation.
- ▶ Compiler Optimization Level "O1" only, owing to complexity of Mapping Problem.

Test Setup


For testing, ARM Cortex A5 core was used. Results from ARM Performance Measurement Unit were compared.

- ▶ Number of Cache Misses.
- ▶ Total Cycles.

Accurate Prediction of Cache Misses, will verify the correctness of the instrumentation.


ADPCM Benchmark

	HCS	Actual	Accuracy
Total Cache Miss	91452	91604	99.99%
Total Cycles	46110394	45594862	99.98%



Sieve Benchmark

	HCS	Actual	Accuracy
Total Cache Miss	933109	934250	99.88%
Total Cycles	99068687	100403731	98.67%





S. Chakravarty, Zhuoran Zhao, and A. Gerstlauer.

Automated, retargetable back-annotation for host compiled performance and power modeling.

In *Hardware/Software Codesign and System Synthesis (CODES+ISSS)*, 2013 International Conference on, pages 1–10, Sept 2013.



Kun Lu, D. Muller-Gritschneider, and U. Schlichtmann.

Memory access reconstruction based on memory allocation mechanism for source-level simulation of embedded software.

In *Design Automation Conference (ASP-DAC)*, 2013 18th Asia and South Pacific, pages 729–734, Jan 2013.